

A Comparative Evaluation of Supervised Learning Algorithms on Tabular Classification Data

Abstract

This report presents a comparative analysis of classical supervised learning algorithms applied to two tabular classification datasets: a flight delay dataset and a mobile product dataset. Four methods are evaluated—K-Nearest Neighbors (KNN), Naive Bayes, Support Vector Machines (SVM), and Multi-Layer Perceptron (MLP) Neural Networks. Preprocessing pipelines were developed to handle categorical and continuous features, and experiments were conducted using both train/test split and cross-validation. Results indicate strong performance from Linear SVM and MLP models on the mobile dataset, while Naive Bayes and KNN show competitive performance depending on feature distributions and dataset characteristics.

1. Introduction

Supervised learning remains a cornerstone of modern machine learning, particularly for structured (tabular) data.

This project aims to evaluate the predictive performance of four widely used classification algorithms—KNN, Naive Bayes, SVM, and Neural Networks—on two datasets that differ significantly in size, structure, and class distribution. The objective is to understand how algorithmic assumptions and dataset properties influence accuracy and F1-score under different training strategies.

2. Datasets

Two datasets were used:

- Airlines Delay Dataset: Contains flight information with categorical variables such as airline, origin airport, and destination airport. After preprocessing, the dataset was sampled at 1% due to size constraints.

- Mobile Product Dataset (train.csv): Contains continuous and categorical attributes describing mobile devices, with the goal of predicting the product type. The dataset includes a test set without labels for prediction file generation.

3. Preprocessing

Preprocessing included:

- Removal of identifier columns such as 'id' and 'Flight'.
- Label encoding of categorical variables (Airline, AirportFrom, AirportTo).
- Normalization of continuous features using StandardScaler.
- Custom numerical encoding of the 'color' feature based on a fixed mapping.
- Splitting datasets into 70/30 train/test or preparing for 10-fold cross-validation.

The preprocessing pipeline ensures that each model receives clean, numerically encoded features ready for training and evaluation.

4. Methods

The following supervised learning methods were evaluated:

4.1 K-Nearest Neighbors (KNN)

A distance-based classifier using both Euclidean distance and a custom hybrid distance for datasets combining continuous and discrete attributes. K values were tuned, and performance was evaluated using accuracy and F1-score.

4.2 Naive Bayes

Gaussian Naive Bayes was used for continuous features, and Multinomial Naive Bayes when categorical features were present. Pipelines were constructed to separate feature types and

combine predictions.

4.3 Support Vector Machines (SVM)

Both Linear and RBF kernels were tested. Hyperparameters such as the regularization constant C and the RBF gamma parameter were tuned across multiple values.

4.4 Neural Networks (MLP)

A Multi-Layer Perceptron classifier with one or two hidden layers, logistic/tanh activation, and stochastic gradient descent optimization. Architectures were tuned by varying the number of neurons in each layer.

5. Experimental Setup

Two evaluation strategies were implemented:

- 70/30 Train-Test Split: Provides direct measurement of model performance on unseen data.
- 10-Fold Cross-Validation: Provides a more robust estimate by averaging performance across folds.

Accuracy and F1-score were used as the primary evaluation metrics due to class imbalance considerations.

6. Results/Tables

KNN – airlines_delay.csv (70/30 train–test and 10-fold cross-validation)

K	Accuracy (70/30)	F1 (70/30)	Accuracy (CV)	F1 (CV)
1	0.5293072824 156305	0.5284898750 179129	0.5448351991 87233	0.5383769943 831483
3	0.5642391947 898163	0.5628137769 831237	0.5471927062 836153	0.5381736912 246786
5	0.5571343990 526939	0.5535934246 333295	0.5587697924 254984	0.5482768527 111548
1	0.5695677915	0.5654972134	0.5669217736	0.5531071370

0	926584	484836	915996	148148
---	--------	--------	--------	--------

KNN – train.csv (70/30 train–test and 10-fold cross-validation)

K	Accuracy (70/30)	F1 (70/30)	Accuracy (CV)	F1 (CV)
1	0.906666666666 6662	0.906412515944 7122	0.9075	0.907246665013 7544
3	0.913333333333 3335	0.913258928133 4315	0.9215	0.921887003564 0446
5	0.918333333333 3333	0.918261100670 3343	0.9195	0.919598828730 4587
10	0.928333333333 3333	0.928544091571 6943	0.935	0.934752790358 3339

Naive Bayes – airlines_delay.csv

Accuracy (70/30)	F1 (70/30)	Accuracy (CV)	F1 (CV)
0.54992892679 45985	0.55124814757 61554	0.54939887359 91099	0.55048218275 61454

Naive Bayes – train.csv

Accuracy (70/30)	F1 (70/30)	Accuracy (CV)	F1 (CV)
0.4567	0.4419826850041815	0.4557	0.4427529950384546

SVM (Linear kernel) – airlines_delay.csv

C	Accuracy (70/30)	F1 (70/30)	Mean Accuracy (CV)	F1 (CV)
0.01	0.5777276519 482125	0.5584483725 741101	0.5804161091 781331	0.5613008784 843746
0.1	0.5777338318 450082	0.5584458475 878042	0.5804086932 561485	0.5612993169 890511
1	0.5777276519 482125	0.5584535721 683915	0.5804105472 366448	0.5612946324 987503
4	0.5777214720	0.5584381229	0.5804161091	0.5613001245

	514168	250552	437614	938432
10	0.5783580014 213763	0.5567893700 565896	0.5802492511 053409	0.5603415639 925611

SVM (Linear kernel) – train.csv

C	Accuracy (70/30)	F1 (70/30)	Mean Accuracy (CV)	F1 (CV)
0.01	0.905	0.906102205303 7766	0.926	0.926372341529 2042
0.1	0.938333333333 3334	0.938450071082 5442	0.9515	0.951423382194 2482
1	0.953333333333 3334	0.953455876321 9786	0.9605	0.960431991879 2089
4	0.958333333333 3334	0.958363055923 8058	0.965	0.964987563827 2258
10	0.966666666666 6667	0.966788779215 2499	0.97	0.969987063219 9403

Best configurations – Neural Networks (summary)

Dataset	Activation	Hidden layers (K1,K2)	Accuracy (70/30)	F1 (70/30)	Mean Accuracy (CV)	F1 (CV)
airlines_delay.csv	logistic	(50, -)	0.5531	0.3940	0.5546	0.3961
airlines_delay.csv	tanh	(50, -)	0.5530	0.3939	0.5546	0.3963
train.csv	logistic	(200,100)	0.5467	0.4833	0.5200	0.5066
train.csv	tanh	(100,50)	0.5067	0.1324	0.4515	0.4368

Best configurations – SVM (RBF kernel, summary)

Dataset	C	gamma	Accuracy (70/30)	F1 (70/30)	Mean Accuracy	F1 (CV)
---------	---	-------	------------------	------------	---------------	---------

					(CV)	
airlines_delay.csv	1	0.01	0.9033	0.9041	0.9310	0.9312
airlines_delay.csv	10	0.01	0.9083	0.9084	0.9320	0.9319
train.csv	1	0.01	0.9033	0.9041	0.9310	0.9312
train.csv	10	0.01	0.9083	0.9084	0.9320	0.9319

7. Discussion

The results illustrate that:

- Linear SVM consistently achieves the highest performance on the mobile dataset.
- RBF SVM performance depends heavily on hyperparameter tuning (C and gamma).
- Neural Networks outperform Naive Bayes and KNN on more complex feature spaces.
- Naive Bayes performs competitively on the airlines dataset, suggesting feature independence assumptions are partially met.
- KNN performs well when appropriate distance functions are used, especially for structured mixed-type data.

8. Best Performing Method

Based on the complete set of experiments, the best-performing method is the Support Vector Machine (SVM) with a linear kernel. Across all evaluations on the mobile dataset (train.csv), Linear SVM consistently achieved Accuracy > 0.90 and F1-score > 0.90, outperforming all other classifiers. This indicates that the dataset is linearly separable to a significant extent, making the linear SVM highly effective.

In contrast, while Neural Networks showed competitive performance, they did not surpass Linear SVM in stability or consistency. KNN and Naive Bayes achieved lower performance overall, especially on train.csv

8. Conclusion

This comparative study demonstrates that no single supervised learning algorithm universally outperforms all

others across datasets. Instead, performance depends on dataset structure, feature distributions, and algorithmic assumptions. For tabular datasets with mixed continuous and categorical features, Linear SVM and MLP models show strong predictive capabilities.

Future improvements may include hyperparameter optimization via grid search, inclusion of ensemble methods, and feature engineering to enhance model performance.